

Hybrid Three-Dimensional LTD/CFF Representations

Implementation and theory note

April 2026

Abstract

This note describes the derivative-free three-dimensional representation implemented in `hybrid3d.py`. The construction combines loop-tree duality (LTD) with the causal/CFF representation in order to handle repeated propagators while keeping the numerator as a sampled black box. It also describes the bounded-energy-degree extension used when edge-momentum representation (EMR) numerators contain powers above the affine CFF class.

The document is organized as follows. Chapter 1 gives the mathematical construction, with all dependencies made explicit. Chapter 2 explains how the Python implementation and the oriented JSON schema encode the formulae. Chapter 3 reports current Symbolica evaluator timings from the example scripts.

Contents

1	Theoretical Construction	2
1.1	Framing, Notation, LTD, and CFF	2
1.2	Repeated Propagators Without Numerator Derivatives	3
1.2.a	The Fully Generic Formula	3
1.2.b	The Box with a Triple Repeated Propagator	7
1.3	Higher Energy Powers in Numerators	9
1.3.1.a	Quadratic Bounds	9
1.3.1.b	One-Loop Box with One Quadratic Edge	10
1.3.2.a	Bounds Above Quadratic	11
1.3.2.b	One-Loop Box with One Cubic Edge	13
2	Python Implementation and JSON Format	14
2.1	Module Structure	14
2.2	Oriented JSON Semantics	14
2.3	Surface Cache	15
2.4	Orientation Labels and Variants	15
2.5	Command-Line Workflow	15
2.6	Symbolica Evaluators	15
3	Performance	17
3.1	Benchmark Setup	17
3.2	Repeated Box	17
3.3	One-Loop 10-Gon	17
3.4	Five-Loop Four-External Benchmarks	18
3.5	Precision Stability	19

Chapter 1

Theoretical Construction

1.1 Framing, Notation, LTD, and CFF

Let $\mathcal{G}$ be a connected scalar graph with L loop momenta and internal edge set E . A loop-momentum basis (LMB) is denoted by $\mathcal{B} = (k_1, \dots, k_L)$. For an edge $e \in E$,

$$q_e(k, p) = \sum_{r=1}^L \rho_{er} k_r + \sum_{a=1}^{n_{\text{ext}}} \eta_{ea} p_a, \quad \rho_{er}, \eta_{ea} \in \mathbb{Z}, \quad (1.1)$$

and

$$q_e^0(k^0, p^0) = \rho_e \cdot k^0 + \eta_e \cdot p^0, \quad \omega_e(\vec{k}, p, m) = \sqrt{\vec{q}_e^2 + m_e^2}. \quad (1.2)$$

The Feynman denominator is

$$D_e = (q_e^0)^2 - \omega_e^2 + i0. \quad (1.3)$$

For a numerator N in EMR variables, the D -dimensional input integral is

$$I_{\mathcal{G}, \mathcal{B}}[N] = \int_{\mathbb{R}^{4L}} \prod_{r=1}^L \frac{d^4 k_r}{(2\pi)^4} N(\{q_e(k, p)\}_{e \in E}) \prod_{e \in E} D_e^{-a_e}. \quad (1.4)$$

The powers a_e are normally one; repeated propagators are encoded below as explicit repeated edges or, equivalently, as collapsed channels with powers.

LTD dependency data. The LTD construction is an algorithmic residue expansion of (1.4) after choosing

- the graph $\mathcal{G}$,
- the loop-momentum basis $\mathcal{B}$,
- a contour-closure prescription κ , i.e. the side on which each loop-energy integration is closed.

We denote the resulting spatial integrand by

$$I_{\text{LTD}}^{(3)}[\mathcal{G}, \mathcal{B}, \kappa; N](\vec{k}; p, m) = \sum_{\mathfrak{o} \in \mathcal{O}_{\text{LTD}}(\mathcal{G}, \mathcal{B}, \kappa)} N\left(Q_{\mathfrak{o}}^{\text{LTD}}(\vec{k}; p, m)\right) W_{\mathfrak{o}}^{\text{LTD}}(\vec{k}; p, m). \quad (1.5)$$

The map $Q_{\mathfrak{o}}^{\text{LTD}}$ gives every EMR edge energy after solving the cut equations. The weight $W_{\mathfrak{o}}^{\text{LTD}}$ is a product of cut half-edge factors and LTD dual surfaces. In this note we do not rederive LTD; we use the multiloop algorithmic construction of Ref. [1].

CFF dependency data. For CFF/causal representations we use the graph as the primary input and assume the canonical source/sink contraction ordering used by the implementation. Thus

$$I_{\text{CFF}}^{(3)}[\mathcal{G}; N](\vec{k}; p, m) = \sum_{\mathfrak{o} \in \mathcal{O}_{\text{CFF}}(\mathcal{G})} N\left(Q_{\mathfrak{o}}^{\text{CFF}}(\vec{k}; p, m)\right) W_{\mathfrak{o}}^{\text{CFF}}(\vec{k}; p, m). \quad (1.6)$$

The CFF weights use causal E -surfaces in the denominator. The construction is the one of Refs. [2, 3]. The standard black-box numerator assumption is that N is at most affine in each EMR edge energy q_e^0 .

Surface notation. For any integer vector $c_e \in \{-1, 0, 1\}$ and external-energy shift Δ , define the signed causal surface

$$\eta(c, \Delta) = \sum_e c_e \omega_e + \Delta. \quad (1.7)$$

An E -surface has all non-zero internal-energy coefficients with the same orientation after the canonical CFF construction. An H -surface is a more general signed hyperboloid surface. Pure CFF denominator factors use only E -surfaces; LTD and hybrid terms may use H -surfaces when the residue geometry requires them. Numerator-side surfaces are allowed in the numerator and do not introduce denominator singularities.

1.2 Repeated Propagators Without Numerator Derivatives

Repeated propagators correspond to groups

$$R_g = \{r_{g,1}, \dots, r_{g,\nu_g}\}, \quad \nu_g \geq 2, \quad (1.8)$$

whose members have the same canonical routing and mass, possibly with reversed edge orientation. Let Q_g^0 be the canonical repeated-channel energy and ω_g the common on-shell energy. The purpose of the hybrid construction is to evaluate the repeated-channel residues without applying derivatives to the black-box numerator.

1.2.a The Fully Generic Formula

Start from the four-dimensional integral split into non-repeated edges S and repeated groups G_{rep} :

$$\begin{aligned} I_{\mathcal{G}, \mathcal{B}}[N] &= \int \prod_{r=1}^L \frac{d^4 k_r}{(2\pi)^4} N(\{q_e(k, p)\}) \prod_{s \in S} D_s^{-1} \prod_{g \in G_{\text{rep}}} \prod_{a=1}^{\nu_g} D_{r_{g,a}}^{-1} \\ &= \int \prod_{r=1}^L \frac{d^4 k_r}{(2\pi)^4} N(\{q_e(k, p)\}) \prod_{s \in S} D_s^{-1} \prod_{g \in G_{\text{rep}}} D_g^{-\nu_g}. \end{aligned} \quad (1.9)$$

The second equality is only a denominator identity. The numerator still has one EMR argument for every original edge $e \in E$, including all repeated copies. Thus the collapsed channel g has a single denominator energy Q_g^0 , but the numerator still sees $q_{r_{g,a}}^0$ for $a = 1, \dots, \nu_g$.

Let $\bar{E} = S \sqcup G_{\text{rep}}$ be the collapsed edge set and assign powers

$$n_{\bar{e}} = \begin{cases} 1, & \bar{e} = s \in S, \\ \nu_g, & \bar{e} = g \in G_{\text{rep}}. \end{cases}$$

For every rank- L collapsed cut basis $B \subset \bar{E}$ and pole-sign vector σ_B , solve the cut equations

$$Q_b^0(k^0, p^0) = \sigma_b \omega_b, \quad b \in B, \quad (1.10)$$

for the loop energies:

$$k^0 = K_{B,\sigma_B}(\vec{k}; p, m). \quad (1.11)$$

The Jacobian associated with this change of variables is

$$J_B = \det(\rho_{br})_{b \in B, r=1, \dots, L}. \quad (1.12)$$

With the contour convention used in the implementation, the sign carried by the corresponding simple LTD residue is denoted by $s_{B,\sigma_B} \in \{\pm 1\}$. The direct residue theorem gives the strict equality

$$I_{\mathcal{G},B}[N] = \int \prod_{r=1}^L \frac{d^3 k_r}{(2\pi)^3} \sum_{B,\sigma_B} s_{B,\sigma_B} \left[\prod_{b \in B} \frac{1}{(n_b - 1)!} \left(\frac{\partial}{\partial Q_b^0} \right)^{n_b-1} \frac{N(\{q_e^0(k^0, p^0), \vec{q}_e\})}{\prod_{b \in B} (Q_b^0 + \sigma_b \omega_b)^{n_b} \prod_{\bar{e} \notin B} D_{\bar{e}}^{n_{\bar{e}}}} \right]_{k^0=K_{B,\sigma_B}}. \quad (1.13)$$

The determinant J_B is already included in the LTD residue sign and normalization used to define s_{B,σ_B} . For the DOT examples all non-zero J_B are ± 1 . If a different integer LMB gives $|J_B| \neq 1$, the residue weight is multiplied by $1/|J_B|$.

Equation (1.13) is not yet acceptable when a powered channel is cut, because the derivatives act on N . Define the set of active original edges

$$\mathcal{A}_B = \{e \in E : q_e^0 \text{ depends on } Q_b^0 \text{ for some } b \in B \text{ with } n_b > 1\}. \quad (1.14)$$

For the affine EMR class, one exact algebraic route is to insert the multilinear pole interpolation identity

$$N(\{q_e^0\}) = \sum_{\theta \in \{\pm 1\}^{\mathcal{A}_B}} N(\{q_{e,\theta}^0\}) \prod_{e \in \mathcal{A}_B} \frac{1 + \theta_e q_e^0 / \omega_e}{2}, \quad (1.15)$$

with $q_{e,\theta}^0 = \theta_e \omega_e$ for active edges and the ordinary affine value for inactive edges. Since N is affine in each active EMR energy, differentiating (1.15) differentiates only the known interpolation factors and never the black-box function values.

The form actually used for the hybrid repeated sectors is the equivalent physical CFF-localized form. To derive it independently, temporarily split the repeated denominator into distinct on-shell energies $\omega_{g,a}$, $a = 1, \dots, \nu_g$, and use the linear-pole decomposition

$$\frac{1}{(Q_g^0)^2 - \omega_{g,a}^2 + i0} = \sum_{\tau_{g,a}=\pm 1} \frac{\tau_{g,a}}{2\omega_{g,a}} \frac{1}{Q_g^0 - \tau_{g,a}\omega_{g,a} + i0\tau_{g,a}}. \quad (1.16)$$

Each linear pole has the causal exponential representation

$$\frac{1}{Q_g^0 - \tau\omega + i0\tau} = -i\tau \int_0^\infty d\lambda \exp[i\tau\lambda Q_g^0 - i\lambda\omega]. \quad (1.17)$$

For a repeated group this gives a sign orbit $\tau_g = (\tau_{g,1}, \dots, \tau_{g,\nu_g})$ and the phase

$$\Phi_{\tau_g} = \exp \left[iQ_g^0 \sum_{a=1}^{\nu_g} \tau_{g,a} \lambda_{g,a} - i \sum_{a=1}^{\nu_g} \lambda_{g,a} \omega_{g,a} \right], \quad \lambda_{g,a} \geq 0. \quad (1.18)$$

The loop-energy integrations are now elementary Fourier integrations. For a complete set of CFF/LTD pole choices P , with signs τ_j , energies ω_j , and loop-energy rows ρ_j , one obtains the localizing constraint

$$\prod_{r=1}^L 2\pi \delta \left(\sum_{j \in P} \tau_j \lambda_j \rho_{jr} \right). \quad (1.19)$$

Solving (1.19) with $\lambda_j \geq 0$ is the standard CFF source/sink contraction step: the surviving sign assignments are precisely the causal orientations. In a repeated group the individual $\tau_{g,a}$ survive as numerator-sampling signs, while the localized collapsed chain has a single sign $\chi_g = \pm 1$ for the energy Q_g^0 . If all repeated-copy signs are equal, χ_g is forced to that common sign. If the orbit is mixed, both $\chi_g = \pm 1$ are possible localizations. This is exactly the rule implemented as

$$\tau_g = (+, +, +) \Rightarrow \chi_g = +, \quad \tau_g = (-, -, -) \Rightarrow \chi_g = -, \quad \text{mixed } \tau_g \Rightarrow \chi_g = \pm.$$

Only after solving the delta constraints do we take the equal-energy limit $\omega_{g,a} \rightarrow \omega_g$. The split expression is an ordinary simple-pole CFF/LTD expression, hence the limit is equal to the confluent residue (1.13). The difference is organizational:

- the confluent derivative form differentiates the known denominator factors and would differentiate N unless (1.15) is inserted first;
- the physical CFF-localized form samples N directly at repeated-copy on-shell sign orbits (τ_g, χ_g) , then combines those samples with rational affine weights.

For an affine EMR numerator the two forms are identical because both are the same equal-energy limit of the split simple-pole identity. They differ only in the basis of numerator calls: the direct interpolation basis uses the formal active edge signs θ_e , while the physical basis uses the CFF-localized orbits (τ_g, χ_g) . The latter is the one printed by the current hybrid JSON labels such as `physG3tmmcm`.

It remains to define the rational weights multiplying those physical samples. Let $\mathcal{X}_{B,\sigma}$ be the finite set of physical samples obtained from the above sign-orbit rule for the active repeated groups of the cut (B, σ_B) . A sample $x \in \mathcal{X}_{B,\sigma}$ determines complete loop and edge energy maps,

$$K_x^0(\vec{k}; p, m), \quad q_{e,x}^0(\vec{k}; p, m).$$

Write every component of every map as an affine vector in the symbols

$$1, \quad p_a^0, \quad \omega_{\bar{e}},$$

where repeated copies share the collapsed symbol ω_g . Let R_x be the column vector containing all those affine coefficients for sample x . For a numerator-derivative multi-index β , let T_β be the target affine vector:

- for $|\beta| = 0$, T_0 is the ordinary cut map $K_{B,\sigma_B}^0, q_{e,B,\sigma_B}^0$;
- for $|\beta| = 1$, T_β is the first derivative of those maps with respect to the corresponding repeated cut energy, multiplied by $2\omega_g$. The compensating factor $1/(2\omega_g)$ is stored as one additional repeated half-edge in the JSON.

The sample coefficients $\lambda_{\beta x}$ are defined by the affine system

$$\sum_{x \in \mathcal{X}_{B,\sigma}} \lambda_{\beta x} = \delta_{\beta 0}, \quad \sum_{x \in \mathcal{X}_{B,\sigma}} \lambda_{\beta x} R_x = T_\beta. \quad (1.20)$$

If the system is underdetermined, the implementation uses the minimum Euclidean norm solution. Equivalently, after discarding dependent rows, if R is the matrix with columns R_x and t_β is the target vector including the constant row, then

$$\lambda_\beta = R^T (R R^T)^{-1} t_\beta. \quad (1.21)$$

This equation is not a fit or approximation: the affine support of $\mathcal{X}_{B,\sigma}$ is checked exactly over rational numbers.

Substituting (1.15) into (1.13) gives

$$I_{\mathcal{G},\mathcal{B}}[N] = \int \prod_{r=1}^L \frac{d^3 k_r}{(2\pi)^3} I_{\mathcal{G},\mathcal{B},\kappa}^{\text{HybridLTD}}[N](\vec{k}; p, m), \quad (1.22)$$

$$I_{\mathcal{G},\mathcal{B},\kappa}^{\text{HybridLTD}}[N] = \sum_{\mathfrak{o} \in \mathcal{O}_H(\mathcal{G},\mathcal{B},\kappa)} N\left(Q_{\mathfrak{o}}^H(\vec{k}; p, m)\right) W_{\mathfrak{o}}^H(\vec{k}; p, m). \quad (1.23)$$

Here one orientation is

$$\mathfrak{o} = (B, \sigma_B, x, \beta, \gamma, \mu),$$

where $x \in \mathcal{X}_{B,\sigma}$, β is the part of the confluent derivative reconstructed from numerator samples, $\gamma = \alpha - \beta$ acts on known denominator factors, and μ labels one algebraic term in the known denominator derivative expansion. Here

$$\alpha_b = n_b - 1, \quad b \in B.$$

The numerator is never differentiated; it is only sampled at

$$Q_{\mathfrak{o}}^H = \{q_{e,x}^0, \vec{q}_e\}_{e \in E}.$$

The denominator derivative coefficient is also explicit. Each known factor is of the form

$$F_j(y) = \zeta_j(y)^{-p_j}, \quad \zeta_j(y) = \zeta_j(0) + a_j \cdot y,$$

where $y = (Q_b^0 - \sigma_b \omega_b)_{b \in B}$, p_j is either a propagator power or a surface power, and a_j is the vector of derivatives with respect to the cut-basis energies. A multi-derivative assignment δ_j with $\sum_j \delta_j = \gamma$ contributes

$$d_{\{\delta_j\}} = \frac{\gamma!}{\prod_j \delta_j!} \prod_j (-1)^{|\delta_j|} (p_j)_{|\delta_j|} a_j^{\delta_j}, \quad (p)_r = p(p+1) \cdots (p+r-1). \quad (1.24)$$

For multi-indices, $\gamma! = \prod_b \gamma_b!$, $\delta_j! = \prod_b (\delta_j)_b!$, and

$$a_j^{\delta_j} = \prod_b (a_j)_b^{(\delta_j)_b}.$$

For a surface factor $\zeta_j = \eta_j$, the derivative raises the denominator power of that cached surface. For a cut half-edge factor $(2\omega_b)^{-p_j}$, it raises the number of stored half-edge entries.

The complete rational coefficient multiplying one JSON variant is therefore

$$c_{\mathfrak{o}} = s_{B,\sigma_B} \frac{1}{\alpha!} \binom{\alpha}{\beta} \lambda_{\beta x} d_{\mu}, \quad (1.25)$$

where d_{μ} is the value of (1.24) for the derivative assignment $\mu = \{\delta_j\}$, and

$$\alpha! = \prod_{b \in B} \alpha_b!, \quad \binom{\alpha}{\beta} = \prod_{b \in B} \binom{\alpha_b}{\beta_b}.$$

The weight is the finite product

$$W_{\mathfrak{o}}^H = c_{\mathfrak{o}} \prod_{h \in H_{\mathfrak{o}}} \frac{1}{2\omega_h} \prod_{\eta \in D_{\mathfrak{o}}} \eta(\vec{k}; p, m)^{-m_{\mathfrak{o},\eta}} \prod_{\zeta \in N_{\mathfrak{o}}} \zeta(\vec{k}; p, m). \quad (1.26)$$

Here $H_{\mathfrak{o}}$ is the multiset of half-edge ids, so each $h \in H_{\mathfrak{o}}$ contributes $(2\omega_h)^{-1}$. The denominator multiset $D_{\mathfrak{o}}$ contains cached surfaces

$$\eta(\vec{k}; p, m) = \sum_e c_e \omega_e + \sum_a d_a p_a^0 + c_0,$$

with $c_e, d_a, c_0 \in \mathbb{Z}$. The numerator multiset N_o contains cached numerator-only surfaces of the same linear form; each $\zeta \in N_o$ is such a known linear factor in on-shell and external energies. Thus the previously used symbol $\zeta(\vec{k}; p, m)$ means exactly a cached linear surface factor appearing in the numerator, not an additional denominator. If no repeated channel is present, then $n_b = 1$ for all cuts, $\mathcal{A}_B = \emptyset$, and (1.23) collapses exactly to the LTD representation with the same $(\mathcal{G}, \mathcal{B}, \kappa)$.

1.2.b The Box with a Triple Repeated Propagator

Consider the one-loop box routing

$$q_0 = k, \quad q_1 = k + p_1, \quad q_2 = k + p_1 + p_2, \quad Q = k + p_1 + p_2 + p_3,$$

with three repeated copies of Q . The collapsed denominator is

$$D_0 D_1 D_2 D_Q^3.$$

Write $\Delta_i = p_1^0 + \dots + p_i^0$ and define the causal surfaces

$$E_{ij}^{\sigma_i \sigma_j} = \sigma_i \omega_i + \sigma_j \omega_j + \Delta_i - \Delta_j, \quad (1.27)$$

$$E_{iQ}^{\sigma_i \sigma_Q} = \sigma_i \omega_i + \sigma_Q \omega_Q + \Delta_i - \Delta_3. \quad (1.28)$$

The three simple-cut sectors are

$$I_i^{\text{simple}} = - \frac{N(q_0^0, q_1^0, q_2^0, Q^0, Q^0, Q^0)_{k^0 = -\Delta_i + \omega_i}}{2\omega_i E_{ij}^{++} E_{i\ell}^{++} (E_{iQ}^{++})^3}, \quad \{i, j, \ell\} = \{0, 1, 2\}. \quad (1.29)$$

The signs shown correspond to the upper on-shell pole convention used in the examples; the opposite closure is obtained by replacing every displayed $+\omega$ by the corresponding signed on-shell energy.

For the repeated cut used in the printed example, the collapsed LTD cut is $Q^0 = +\omega_Q$, hence

$$k^0 = \omega_Q - \Delta_3.$$

The six denominator surfaces generated by the three non-cut edges are

$$\eta_{18} = -\omega_0 + \omega_Q - \Delta_3, \quad \eta_{19} = \omega_0 + \omega_Q - \Delta_3, \quad (1.30)$$

$$\eta_{20} = -\omega_1 + \omega_Q - \Delta_2, \quad \eta_{21} = \omega_1 + \omega_Q - \Delta_2, \quad (1.31)$$

$$\eta_{22} = -\omega_2 + \omega_Q - \Delta_1, \quad \eta_{23} = \omega_2 + \omega_Q - \Delta_1. \quad (1.32)$$

These are precisely surface ids 18, ..., 23 in the CLI output. The overview table prints a compact subset of denominator ids, but `-show-details-for-orientation` displays the full factorization tree containing all six surfaces.

Let $\tau = (\tau_3, \tau_4, \tau_5)$ be the physical repeated-copy signs and χ the localized collapsed-chain sign. A physical numerator call is

$$N_{\tau, \chi} = N(\chi \omega_Q - \Delta_3, \chi \omega_Q - \Delta_2, \chi \omega_Q - \Delta_1, \tau_3 \omega_Q, \tau_4 \omega_Q, \tau_5 \omega_Q). \quad (1.33)$$

The admissible physical sample set is

$$\mathcal{X}_Q = \{(+++, +), (---, -)\} \cup \{(\tau, +), (\tau, -) : \tau \in \{\pm\}^3, \tau \neq +++, \tau \neq ---\}. \quad (1.34)$$

For each $x = (\tau, \chi) \in \mathcal{X}_Q$, the affine coefficients λ_{0x} and λ_{1x} defined by (1.20) reconstruct respectively the numerator itself and the first numerator derivative. The full table is

sample $x = (\tau, \chi)$	λ_{0x}	λ_{1x}
(+ + +, +)	43/112	5/8
(- - -, -)	-27/112	-5/8
(- - +, -)	-3/56	-1/4
(- - +, +)	1/112	-1/8
(- + -, -)	-3/56	-1/4
(- + -, +)	1/112	-1/8
(- + +, -)	15/112	1/8
(- + +, +)	11/56	1/4
(+ - -, -)	-3/56	-1/4
(+ - -, +)	1/112	-1/8
(+ - +, -)	15/112	1/8
(+ - +, +)	11/56	1/4
(+ + -, -)	15/112	1/8
(+ + -, +)	11/56	1/4

The first column sums to one and reconstructs the base affine numerator map. The second column sums to zero and reconstructs the derivative map, with one extra factor $1/(2\omega_Q)$.

For the repeated box, $\alpha = 2$. Hence $\beta = 0, \gamma = 2$ or $\beta = 1, \gamma = 1$. Define

$$P = \prod_{i=18}^{23} \eta_i^{-1}, \quad h_r = \sum_{18 \leq i_1 \leq \dots \leq i_r \leq 23} \eta_{i_1}^{-1} \dots \eta_{i_r}^{-1}, \quad h_0 = 1.$$

The denominator derivative terms are

$$\gamma = 2: \quad 2 \frac{Ph_2}{(2\omega_Q)^3} + 6 \frac{Ph_1}{(2\omega_Q)^4} + 12 \frac{P}{(2\omega_Q)^5}, \quad (1.35)$$

$$\gamma = 1: \quad - \frac{Ph_1}{(2\omega_Q)^3} - 3 \frac{P}{(2\omega_Q)^4}. \quad (1.36)$$

These equations are exactly the two factorization trees printed in the JSON: Ph_2 is the large tree with all weakly ordered pairs of surfaces, Ph_1 is the smaller tree, and P is the single chain [18, 19, 20, 21, 22, 23].

The complete repeated-cut contribution is therefore

$$I_Q^{\text{rep}} = \sum_{x \in \mathcal{X}_Q} N_x \left[-\frac{\lambda_{0x}}{2} \left(2 \frac{Ph_2}{(2\omega_Q)^3} + 6 \frac{Ph_1}{(2\omega_Q)^4} + 12 \frac{P}{(2\omega_Q)^5} \right) - \lambda_{1x} \left(-\frac{Ph_1}{(2\omega_Q)^4} - 3 \frac{P}{(2\omega_Q)^5} \right) \right]. \quad (1.37)$$

The overall minus sign is the one-loop LTD residue sign for the Q -cut in this routing. Equation (1.37) contains no undefined coefficient: every rational number is either one of the affine weights in the table or one of the derivative factors in (1.35)–(1.36).

For the concrete orientation requested in the CLI output,

$$x = (- - -, -), \quad N_x = N_{---, -}.$$

The coefficients are

$$\lambda_{0x} = -\frac{27}{112}, \quad \lambda_{1x} = -\frac{5}{8}.$$

Using (1.25), the $\beta = 0, \gamma = 2$ variants are

$$-\frac{1}{2} \lambda_{0x} (2, 6, 12) = \left(\frac{27}{112}, \frac{81}{112}, \frac{81}{56} \right),$$

with half-edge powers 3, 4, 5. The $\beta = 1, \gamma = 1$ variants include the extra numerator-derivative half-edge and have

$$-\lambda_{1x}(-1, -3) = \left(-\frac{5}{8}, -\frac{15}{8}\right),$$

with half-edge powers 4, 5. These five numbers are exactly

variant label	prefactor	half-edge multiset
B3 b0 g2 n0:physG3tmmcm	27/112	[3, 3, 3]
B3 b0 g2 n0:physG3tmmcm	81/112	[3, 3, 3, 3]
B3 b0 g2 n0:physG3tmmcm	81/56	[3, 3, 3, 3, 3]
B3 b1 g1 nD0phys:physG3tmmcm	-5/8	[3, 3, 3, 3]
B3 b1 g1 nD0phys:physG3tmmcm	-15/8	[3, 3, 3, 3, 3].

The label **tmmcm** means $\tau = (-, -, -)$ and $\chi = -$; **b0/g2** means $\beta = 0, \gamma = 2$, and **b1/g1** means $\beta = 1, \gamma = 1$. The repeated edge signs in the orientation are all minus because the numerator sample uses the physical repeated-copy signs $\tau = (-, -, -)$, even though the denominator tree is the derivative expansion around the collapsed $Q^0 = +\omega_Q$ cut.

The final boxed formula is

$$I_{\square(3)}^{\text{HybridLTD}} = I_0^{\text{simple}} + I_1^{\text{simple}} + I_2^{\text{simple}} + I_Q^{\text{rep}}.$$

1.3 Higher Energy Powers in Numerators

The CFF construction is naturally compatible with numerators affine in each EMR edge energy. When a user supplies a bound d_e for the highest power of each edge energy in N , the implementation first checks loop-energy UV convergence. In every loop-energy direction v , the degree of the numerator minus twice the number of denominators depending on that direction must be strictly smaller than -1 . If this check fails, finite pole data are not enough and the command refuses to build a representation.

1.3.1.a Quadratic Bounds

Let $x = q_e^0$, $\omega = \omega_e$, and suppose N has degree at most two in x , while all other EMR edge energies are held fixed. There is an exact polynomial division

$$N(x) = R_e N(x) + (x^2 - \omega^2) C_e N, \quad (1.38)$$

where the pole-safe affine remainder is

$$R_e N(x) = \frac{x + \omega}{2\omega} N(+\omega) + \frac{\omega - x}{2\omega} N(-\omega), \quad (1.39)$$

and the contact quotient is

$$C_e N = \frac{N(+\omega) - 2N(0) + N(-\omega)}{2\omega^2}. \quad (1.40)$$

Therefore

$$\frac{N}{D_e \prod_{f \neq e} D_f} = \frac{R_e N}{D_e \prod_{f \neq e} D_f} + \frac{C_e N}{\prod_{f \neq e} D_f}. \quad (1.41)$$

The first term is affine in q_e^0 and can be handled by ordinary CFF. The second term is a pinched graph with edge e removed, also handled by CFF. Consequently the pure CFF denominator remains E -surface only. The non-trivial part is to lift the CFF expression of the deleted graph back to an orientation of the original graph. The lift is completely mechanical:

1. Delete the pinched edge set P from the denominator graph and build a CFF representation of the lower graph $\mathcal{G} \setminus P$. If the lower graph factorizes into independent loop-energy components, build one CFF factor per component and multiply the component representations.
2. Every lower-sector CFF term fixes loop energies $K_t^{0, \mathcal{G} \setminus P}$ in the original LMB. If component CFF factors were built separately, the global loop-energy map is reconstructed by solving the full-rank linear system defined by the union of component basis edges.
3. For every original edge, including each deleted edge, compute the lifted EMR energy

$$\hat{q}_{e,t}^0 = \rho_e \cdot K_t^{0, \mathcal{G} \setminus P} + \eta_e \cdot p^0. \quad (1.42)$$

This number is used only for known polynomial factors. It is not automatically the numerator sample value of a deleted edge.

4. The numerator sample value of a deleted quadratic edge is one of $+\omega_e, 0, -\omega_e$, because $C_e N$ in (1.40) depends on those three black-box samples. The original JSON orientation therefore assigns $q_e^0 = +\omega_e, 0$, or $-\omega_e$ to the missing edge and records the corresponding orientation character $+, 0$, or $-$.
5. Any lifted factor $\hat{q}_{e,t}^0$ that remains outside the black-box numerator is stored as a numerator-side surface. In the strictly quadratic contact quotient no such q -factor is needed; for cubic and higher bounds it is needed.
6. The lower-sector CFF denominator surfaces are copied into the original surface cache after replacing every lower-sector edge id by its original edge id. No denominator surface involving a deleted edge is introduced.

The rational coefficients in the JSON follow directly from (1.40). Since half-edge lists encode $(2\omega_e)^{-1}$, the factor $1/(2\omega_e^2)$ is represented as

$$\frac{1}{2\omega_e^2} = \frac{2}{(2\omega_e)^2}.$$

The three samples $N(+\omega_e), N(0), N(-\omega_e)$ therefore carry raw half-edge prefactors $2, -4, 2$. A one-edge CFF pinch contributes one additional sign (-1) in the implementation's CFF orientation convention, so the printed pinched quadratic prefactors are $-2, +4, -2$.

In hybrid mode, the same division is applied to non-affine edge energies before or inside the repeated-channel residue expansion. For non-repeated edges this is identical to the pure CFF lift above. For repeated edges the division is combined with the repeated-channel sample set $\mathcal{X}_{B,\sigma}$. The black-box numerator is still sampled only at the explicit finite set of energies supplied by the JSON orientation map.

1.3.1.b One-Loop Box with One Quadratic Edge

For a non-repeated one-loop box with denominators $D_0 D_1 D_2 D_3$, assume the only non-affine dependence is $d_0 = 2$. Use the routing

$$q_0 = k, \quad q_1 = k + p_1, \quad q_2 = k + p_1 + p_2, \quad q_3 = k + p_1 + p_2 + p_3,$$

and set $\Delta_i = p_1^0 + \dots + p_i^0$. Equations (1.38)–(1.41) give

$$I_{\square}^{\text{CFF}}[N] = I_{\text{CFF}}[\square; R_0 N] + I_{\text{CFF}}[\square/0; C_0 N], \quad (1.43)$$

where $\square/0$ is the triangle obtained by deleting propagator 0. More explicitly,

$$I_{\text{CFF}}[\square; R_0 N] = \sum_{\mathfrak{o} \in \mathcal{O}_{\text{CFF}}(\square)} \left[\frac{q_0^0(\mathfrak{o}) + \omega_0}{2\omega_0} N(q_0^0 = +\omega_0) + \frac{\omega_0 - q_0^0(\mathfrak{o})}{2\omega_0} N(q_0^0 = -\omega_0) \right] W_{\mathfrak{o}}^{\text{CFF}}, \quad (1.44)$$

$$I_{\text{CFF}}[\square/0; C_0 N] = \sum_{\mathfrak{t} \in \mathcal{O}_{\text{CFF}}(\square/0)} \frac{N(+\omega_0) - 2N(0) + N(-\omega_0)}{2\omega_0^2} W_{\mathfrak{t}}^{\text{CFF}}. \quad (1.45)$$

Both $W_{\mathfrak{o}}^{\text{CFF}}$ and $W_{\mathfrak{t}}^{\text{CFF}}$ contain only E -surface denominators.

The lifted triangle terms make the orientation structure concrete. If the triangle CFF term cuts edge $j \in \{1, 2, 3\}$ with sign σ_j , then

$$k^0 = \sigma_j \omega_j - \Delta_j, \quad \widehat{q}_0^0 = \sigma_j \omega_j - \Delta_j.$$

However, in the contact numerator call the deleted edge energy is not $\widehat{q}_0^0$; it is one of the three interpolation samples

$$q_0^0 = +\omega_0, \quad q_0^0 = 0, \quad q_0^0 = -\omega_0.$$

For example, the lower triangle orientation with original signs $(-++)$ on edges 1, 2, 3 contributes three lifted original-box orientations

$$+ - ++, \quad 0 - ++, \quad - - ++,$$

where the first character is the assigned numerator value of the deleted edge. The printed coefficients are

orientation	variant	prefactor	half-edge multiset
$+ - ++$	<code>pinch[0] e0=+</code>	-2	$[0, 0, 1, 2, 3]$
$0 - ++$	<code>pinch[0] e0=0</code>	4	$[0, 0, 1, 2, 3]$
$- - ++$	<code>pinch[0] e0=-</code>	-2	$[0, 0, 1, 2, 3]$.

The denominator surfaces of that lower triangle are the ordinary CFF surfaces among edges 1, 2, 3; in the CLI table they appear, for this orientation, as ids [4, 5]. The two extra half-edge entries 0, 0 encode the factor $(2\omega_0)^{-2}$. No denominator singularity associated with edge 0 remains because D_0 has been cancelled by the contact quotient. This is why the final JSON is still a normal orientation of the original box: all four numerator edge energies are specified, but only the lower triangle denominator surfaces are present.

1.3.2.a Bounds Above Quadratic

For $d_e > 2$, the same idea works but the contact quotient is no longer a constant in the deleted edge energy. Let $a_1, \dots, a_{d_e-1}$ be non-pole interpolation nodes with $a_j \neq \pm 1$, and write

$$u = \frac{q_e^0}{\omega_e}.$$

The implementation uses the deterministic node list

$$0, 2, -2, 3, -3, \dots$$

truncated to $d_e - 1$ nodes. The Lagrange basis polynomial associated with node a_j is

$$L_j(u) = \prod_{\ell \neq j} \frac{u - a_\ell}{a_j - a_\ell}. \quad (1.46)$$

It is characterized by $L_j(a_i) = \delta_{ij}$. The pole remainder is still the affine interpolation at $u = \pm 1$, and the contact quotient is reconstructed from black-box samples by

$$C_e N(q_e^0) = \sum_j \frac{L_j(q_e^0/\omega_e)}{(a_j^2 - 1)\omega_e^2} \left[N(a_j\omega_e) - \frac{a_j + 1}{2} N(+\omega_e) - \frac{1 - a_j}{2} N(-\omega_e) \right]. \quad (1.47)$$

The quotient $C_e N$ is now a known polynomial in the lifted deleted-edge energy $\hat{q}_{e,t}^0$ of (1.42). Expanding each $L_j(u)$ gives

$$L_j(u) = \sum_{r=0}^{d_e-2} \ell_{jr} u^r.$$

A monomial contribution is therefore

$$\ell_{jr} \frac{(\hat{q}_{e,t}^0)^r}{(a_j^2 - 1)\omega_e^{r+2}} \left[N(a_j\omega_e) - \frac{a_j + 1}{2} N(+\omega_e) - \frac{1 - a_j}{2} N(-\omega_e) \right]. \quad (1.48)$$

In JSON this becomes:

- $r + 2$ entries of edge e in `half_edges`, because $\omega_e^{-(r+2)} = 2^{r+2}(2\omega_e)^{-(r+2)}$;
- r numerator-side copies of the lifted surface $\hat{q}_{e,t}^0$, unless that linear form is zero;
- one numerator map with $q_e^0 = a_j\omega_e$, $+\omega_e$, or $-\omega_e$, depending on the black-box sample in the bracket.

For several higher-power edges, the construction is recursive. Pick one active edge, split its contribution into the pole remainder and the contact sector, then repeat the same operation on the remaining bounded edges of the same graph or lower graph. The recursion terminates when every remaining edge is affine. At that point the denominator is evaluated by ordinary CFF or, for a hybrid repeated sector, by the repeated-channel master formula above.

In multiloop lower sectors the deleted graph may decompose into independent loop-energy components. The implementation uses the vector matroid of the remaining loop-energy rows ρ_e : two lower-sector edges are placed in the same component when their rows are connected through a common rank basis. A CFF representation is generated for each full-rank component. The Cartesian product of component orientations fixes enough lower-sector edge energies to solve the original LMB. All component surfaces are then remapped to original edge ids and multiplied. The lifted deleted-edge energy $\hat{q}_e^0$ can be a signed sum or difference of several on-shell energies; when it is a difference, the numerator-side cached factor is an H -surface. This is harmless because it appears only in the numerator. Pure CFF denominator trees remain E -surface only throughout the recursive quadratic construction, including repeated-signature graphs. If a contacted edge belongs to a repeated group, the contact sector simply lowers that repeated denominator power by one copy; the resulting lower graph is again expanded by ordinary CFF minors before any equal-energy limit is taken.

Hybrid mode uses the same bounded-degree algebra. If the high-power edge is not repeated, the lower-sector lift is exactly the one just described. If the high-power edge is repeated, the finite-pole samples $a_j\omega_e$ are folded into the repeated physical sample set and the remaining denominator derivatives are generated by (1.24). The numerator still appears only through explicit sample values.

The pure CFF builder deliberately rejects repeated-signature graphs only when a cap is higher than quadratic. An earlier implementation reused the bounded hybrid confluent limit for the quadratic repeated corner. That expression is a valid hybrid repeated-pole representation, but it can contain H -surfaces in denominator trees and is therefore not the promised E -surface-only pure CFF construction. The implemented quadratic path instead keeps every lower-sector CFF numerator energy map edge by edge and combines equal repeated copies only through CFF E -surface minors.

1.3.2.b One-Loop Box with One Cubic Edge

For the same one-loop box, take $d_0 = 3$. The implementation chooses the two non-pole nodes

$$a_0 = 0, \quad a_1 = 2.$$

The Lagrange basis is

$$L_0(u) = 1 - \frac{u}{2}, \quad L_1(u) = \frac{u}{2}. \quad (1.49)$$

Substituting these two polynomials into (1.47) gives

$$\begin{aligned} C_0 N(q_0^0) = & -\frac{1 - u/2}{\omega_0^2} \left[N(0) - \frac{1}{2}N(+\omega_0) - \frac{1}{2}N(-\omega_0) \right] \\ & + \frac{u/2}{3\omega_0^2} \left[N(2\omega_0) - \frac{3}{2}N(+\omega_0) + \frac{1}{2}N(-\omega_0) \right], \quad u = \frac{\hat{q}_0^0}{\omega_0}. \end{aligned} \quad (1.50)$$

For a lower triangle orientation $\mathbf{t}$, the lifted deleted energy is

$$\hat{q}_{0,\mathbf{t}}^0 = \rho_0 \cdot K_{\mathbf{t}}^{0,\square/0}.$$

For instance, if the triangle cut is edge 1 with sign $-$, then

$$\hat{q}_{0,\mathbf{t}}^0 = -\omega_1 - p_1^0,$$

which appears in the surface cache as a numerator-only E -surface. If the cut is edge 1 with sign $+$, the lifted surface is $\omega_1 - p_1^0$. These are the numerator-only surfaces printed as (e) in the cubic box output.

The monomial coefficients of (1.50) map directly to JSON variants. Before the one-edge CFF pinch sign, the finite-pole contact components for edge 0 are

sample q_0^0	prefactor	half-edge multiset	numerator surface
0	-4	$[0, 0]$	$\emptyset$
0	4	$[0, 0, 0]$	$[\hat{q}_0^0]$
$+\omega_0$	2	$[0, 0]$	$\emptyset$
$+\omega_0$	-4	$[0, 0, 0]$	$[\hat{q}_0^0]$
$-\omega_0$	2	$[0, 0]$	$\emptyset$
$-\omega_0$	-4/3	$[0, 0, 0]$	$[\hat{q}_0^0]$
$2\omega_0$	4/3	$[0, 0, 0]$	$[\hat{q}_0^0]$.

The printed one-loop box CFF table includes the additional pinch sign, so the visible variants are

orientation label	variant	prefactor	meaning
0 - ++	pinch[0] e0=0	4	$-(-4) N(q_0^0 = 0)/(2\omega_0)^2$
0 - ++	pinch[0] e0=0	-4	$-(4) \hat{q}_0^0 N(q_0^0 = 0)/(2\omega_0)^3$
x - ++	pinch[0] e0=2	-4/3	$-(4/3) \hat{q}_0^0 N(q_0^0 = 2\omega_0)/(2\omega_0)^3$
+ - ++	pinch[0] e0=+	-2	$-(2) N(q_0^0 = +\omega_0)/(2\omega_0)^2$
+ - ++	pinch[0] e0=+	4	$-(-4) \hat{q}_0^0 N(q_0^0 = +\omega_0)/(2\omega_0)^3$
- - ++	pinch[0] e0=-	-2	$-(2) N(q_0^0 = -\omega_0)/(2\omega_0)^2$
- - ++	pinch[0] e0=-	4/3	$-(-4/3) \hat{q}_0^0 N(q_0^0 = -\omega_0)/(2\omega_0)^3$.

The label x in $x - ++$ indicates that the deleted edge is sampled at the non-standard finite node $q_0^0 = 2\omega_0$, so its detailed energy map must be read from the orientation details rather than from a single $+$, $-$, or 0 symbol. All denominator surfaces in these rows are the lifted triangle CFF surfaces; $\hat{q}_0^0$ is a numerator-only surface and cannot create a denominator singularity.

Chapter 2

Python Implementation and JSON Format

2.1 Module Structure

The implementation is organized around an evaluator-complete intermediate format.

- `graph_io.py` parses DOT graphs, extracts momentum signatures, detects repeated channels, and builds split-mass reference graphs.
- `graph_signatures.py` implements the `graph_from_signatures` command, which turns products such as `prop(k1+p1,m)` into DOT graphs.
- `orientation_bundle.py` constructs LTD, CFF, hybrid, and bounded-degree contact bundles. This is where repeated channels are collapsed, active numerator maps are enumerated, and CFF lower minors are generated.
- `structure.py` serializes bundles to schema-v6 JSON and contains the arbitrary-precision Python evaluator.
- `pretty.py` renders the JSON in a human-readable form.
- `symbolica_eval.py` compiles a fixed-numerator JSON expression to Symbolica eager and assembly-backed evaluators.

2.2 Oriented JSON Semantics

The top-level JSON contains graph metadata, a surface cache, and a list of orientations. The central invariant is:

one orientation corresponds to one unique EMR edge-energy numerator map.

All denominator terms sharing that numerator call are stored as **variants**. Each variant has its own rational prefactor, half-edge list, numerator-surface list, and factorized denominator tree.

The evaluator uses only the JSON plus the DOT graph. It does not reconstruct a hidden LTD or CFF object. A variant contributes

$$\text{pref} \prod_{e \in \text{half_edges}} (2\omega_e)^{-1} \prod_{s \in \text{num_surfaces}} \eta_s \prod_{t \in \text{denominator tree}} \eta_t^{-1}$$

times the numerator sampled at the orientation's `edge_q0` map.

2.3 Surface Cache

Each surface cache entry records a class, coefficients of internal on-shell energies, and external-energy shifts. The class is printed as **e** or **h**. If a surface appears only in numerator factors, the entry carries `numerator_only=true` and pretty printing displays (**e**) or (**h**). If a numerator-side surface is identical to an existing denominator surface, the same id is reused.

2.4 Orientation Labels and Variants

Compact labels use $+$, $-$, 0 , and x :

- $+$: the numerator samples that edge at $+\omega_e$;
- $-$: the numerator samples that edge at $-\omega_e$;
- 0 : the numerator samples that edge at zero energy;
- x : the edge has a non-trivial affine map, printed by `-show-details`.

The physical origin of a contribution, such as CFF, LTD, pinched contacts, or hybrid basis/interpolation data, is variant metadata rather than part of the orientation identity. This is what keeps the number of numerator calls minimal.

2.5 Command-Line Workflow

The main commands are:

```
python3 hybrid3d.py validate --dot examples/graphs/box_pow3.dot
python3 hybrid3d.py build --family hybrid --dot examples/graphs/box_pow3.dot --pretty
python3 hybrid3d.py test --dot examples/graphs/box_pow3.dot
python3 hybrid3d.py test-cff-ltd --dot examples/graphs/box.dot --energy-degree-bounds
0:2
python3 hybrid3d.py compile --orientation-json demo.json --dot graph.dot --numerator-
expr "1"
python3 hybrid3d.py evaluate --orientation-json demo.json --dot graph.dot --evaluator-
backend symbolica
```

The repeated-propagator three-way test compares pure CFF, hybrid, and the split-mass LTD limiting proxy. The ordinary LTD formula on an unsplit repeated graph is not a valid physics comparison because numerator derivatives are not included there.

2.6 Symbolica Evaluators

The `compile` command accepts one JSON representation and one fixed numerator expression. It reconstructs the JSON expression symbolically, introducing structured functions such as `OSE(qx,qy,qz,m)`, `etaE(i)`, `etaH(i)`, and tagged edge/loop/external momentum components. The resulting Symbolica evaluator is saved both as a shared library and as eager evaluator state. The JSON is updated with relative paths and fingerprints tying the evaluator to the graph, numerator, value type, and orientation structure.

The `evaluate` command can use:

```
builtin,    symbolica_compiled,  symbolica_eager,  symbolica_eager_symjit.
```

Profiling performs ten calls of a batch evaluator. The default batch size is 100, so a profile line reports one thousand repeated samples of the same input kinematics.

The same saved eager evaluator is used by `evaluate -stability`. This mode converts the generated or supplied real kinematic input to decimal precision-tracking numbers with 16 significant digits by default, evaluates with Symbolica's `Evaluator.evaluate_with_prec`, and reports the number of significant digits carried by the final decimal result. The work precision for constants and arithmetic defaults to 80 decimal digits and can be changed with `-stability-work-precision`. Non-finite outputs are assigned zero reported digits. The printed central value is diagnostic only: once the reported digit count is small, the value may have lost enough significance to be numerically meaningless.

Chapter 3

Performance

3.1 Benchmark Setup

The numbers below were generated on April 26, 2026 using the scripts in `examples/scripts`. Each timing is the per-sample time reported by `evaluate -profiling 100`, i.e. ten batch calls of size 100. The 15-external one-loop benchmark is intentionally omitted here: the LTD build is small, but the pure CFF one-loop representation has $2^{15} - 2 = 32766$ sign sectors and the CFF build is not practical for routine documentation runs.

3.2 Repeated Box

The repeated box script `examples/scripts/box_pow3_cli_showcase.sh` compares hybrid and pure CFF for the six-edge graph with a triple repeated propagator. The numerator is `dot(edges[0], ext[0]) + edges[3][0]`.

label	backend	orientations	maps	value type	per sample
hybrid	compiled	17	17	real	$0.0955\ \mu s$
hybrid	eager	17	17	real	$0.521\ \mu s$
hybrid	eager+symjit	17	17	real	$0.104\ \mu s$
hybrid	compiled	17	17	complex	$0.271\ \mu s$
hybrid	eager	17	17	complex	$0.725\ \mu s$
hybrid	eager+symjit	17	17	complex	$0.0989\ \mu s$
CFF	compiled	62	62	real	$0.162\ \mu s$
CFF	eager	62	62	real	$1.44\ \mu s$
CFF	eager+symjit	62	62	real	$0.142\ \mu s$
CFF	compiled	62	62	complex	$0.466\ \mu s$
CFF	eager	62	62	complex	$2.03\ \mu s$
CFF	eager+symjit	62	62	complex	$0.136\ \mu s$

The hybrid representation has fewer numerator maps than CFF, which explains the faster compiled real evaluator. Complex compiled evaluation is slower because the generated function propagates complex arithmetic even though the input kinematics are real. Symjit eager evaluation is competitive for this small expression because the batch is small and the expression has only 17 or 62 maps.

3.3 One-Loop 10-Gon

The script `examples/scripts/one_loop_10_external_runtime_compare.sh` compares numerator-free LTD and CFF on a true one-loop 10-gon.

family	backend	orientations	maps	per sample
LTD	compiled	10	10	$0.234\ \mu s$
LTD	eager	10	10	$0.852\ \mu s$
LTD	eager+symjit	10	10	$0.147\ \mu s$
CFF	compiled	1022	1022	$3.96\ \mu s$
CFF	eager	1022	1022	$36.6\ \mu s$
CFF	eager+symjit	1022	1022	$2.97\ \mu s$

The result is expected: the one-loop CFF orientation count grows like $2^n - 2$, while LTD has only one cut per internal edge. The CFF expression is still fast in absolute terms after compilation, but it is much larger.

3.4 Five-Loop Four-External Benchmarks

The script `examples/scripts/five_loop_symbolica_runtime_compare.sh` compares three numerator choices on a non-repeated five-loop graph and on the repeated five-loop graph. The three numerator choices are 1, one edge-external dot product per internal edge, and the same numerator with three selected edge terms squared.

case	family	maps	compiled	eager	eager+symjit
no repeats, $N = 1$	LTD	137	0.326	1.88	0.250
no repeats, $N = 1$	CFF	930	1.26	12.3	0.892
no repeats, $N = 1$	hybrid	137	0.333	1.87	0.246
no repeats, linear	LTD	137	0.481	3.29	0.379
no repeats, linear	CFF	930	1.96	17.5	1.35
no repeats, linear	hybrid	137	0.484	3.24	0.391
no repeats, squared	LTD	137	0.493	3.45	1.73
no repeats, squared	CFF	4692	18.7	45.9	13.2
no repeats, squared	hybrid	137	0.520	3.31	0.390

case	family	maps	compiled	eager	eager+symjit
repeated, $N = 1$	LTD	137	0.334	1.90	0.272
repeated, $N = 1$	CFF	930	1.26	12.5	0.899
repeated, $N = 1$	hybrid	630	0.823	4.88	0.579
repeated, linear	LTD	137	0.481	3.44	0.408
repeated, linear	CFF	930	1.88	17.4	1.32
repeated, linear	hybrid	630	21.4	37.8	12.4
repeated, squared	LTD	137	0.491	3.53	0.387
repeated, squared	CFF	see text	see text	see text	see text
repeated, squared	hybrid	115	1.84	9.47	1.11

For non-repeated graphs the hybrid family collapses to the LTD structure, so their orientation and map counts agree. CFF is slower because it has more causal sign sectors. The squared CFF no-repeat case is much larger because the bounded-degree contact completion adds pinched CFF minors.

For repeated graphs, the LTD rows are included only as runtime references. The ordinary unsplit LTD expression does not include numerator derivatives for raised propagators; the correct comparison is CFF versus hybrid versus the split-mass LTD limiting test. The repeated unit numerator shows the expected middle ground: hybrid has fewer maps than CFF but more than

LTD. The repeated linear case is slower for hybrid because the confluent interpolation creates heavier denominator variants per numerator map. In the squared repeated case, the pure CFF row now uses the recursive E -surface-only contact/minor construction rather than the older hybrid fallback; the timing table should be regenerated before quoting a stable number for that row. The bounded hybrid reduction has 115 maps for this benchmark.

3.5 Precision Stability

The following tables use `evaluate -stability` with 16 significant digits in the input and 80 decimal digits of work precision. They measure the precision carried by Symbolica’s eager precision-tracking arithmetic through the already compiled structured expression. These measurements are not identity tests; they probe numerical conditioning of the particular representation and kinematic point. When only one or two digits remain, the displayed value should be interpreted as a symptom of cancellation rather than as a trustworthy estimate.

example	family	maps	result digits	tracked value
repeated box	LTD	6	0	∞
repeated box	CFF	62	1	10^{-2}
repeated box	hybrid	17	1	-2×10^3
one-loop 10-gon	LTD	10	2	2.6×10^5
one-loop 10-gon	CFF	1022	1	10^{-1}

For the repeated box, the LTD row is again not a valid repeated-propagator formula; its singular output is a useful reminder that the unsplit LTD structure contains zero repeated-channel surfaces. CFF and hybrid have the same ordinary double value for this point, but the precision-tracking run shows that both are badly conditioned at 16 input digits, with the hybrid central value already unusable. This is consistent with the hybrid expression being a more compact but more cancellation-prone confluent combination.

five-loop case	LTD digits	CFF digits	hybrid digits
no repeats, $N = 1$	2	2	2
no repeats, linear	2	2	2
no repeats, squared	2	2	2
repeated, $N = 1$	0	2	2
repeated, linear	0	2	2
repeated, squared	0	2	2

The five-loop tables show that the tested random point is numerically delicate for all structured sums at ordinary double-like input precision. The non-repeated hybrid rows coincide with LTD, as required by construction. In the repeated graph, the invalid unsplit LTD expression produces NaN and therefore zero digits, while CFF and hybrid retain two tracked digits in all three numerator scenarios. The result supports using the stability mode as a representation-level diagnostic: it highlights where algebraically equivalent JSON structures differ in numerical conditioning even when the ordinary double-precision evaluator returns matching values.

Bibliography

- [1] J. J. Aguilera-Verdugo et al., algorithmic multiloop LTD construction, arXiv:1906.06138.
- [2] Z. Capatti et al., causal/CFF representation construction, arXiv:2211.09653.
- [3] Z. Capatti et al., further CFF representation and box examples, arXiv:2311.14374.
- [4] cLTD and gammaLoop implementations, <https://github.com/alpha100p/cLTD> and <https://github.com/alpha100p/gammaloop>.